

chap01ex

December 19, 2021

1 Examples and Exercises from Think Stats, 2nd Edition

<http://thinkstats2.com>

Copyright 2016 Allen B. Downey

MIT License: <https://opensource.org/licenses/MIT>

```
[31]: from __future__ import print_function, division

import nsfg
import pandas as pd
```

1.1 Examples from Chapter 1

Read NSFG data into a Pandas DataFrame.

```
[51]: preg = nsfg.ReadFemPreg()
preg.head()
```

```
[51]:
```

	caseid	pregordr	howpreg_n	howpreg_p	moscurrp	nowprgdk	pregend1	\
0	1	1	NaN	NaN	NaN	NaN	6.0	
1	1	2	NaN	NaN	NaN	NaN	6.0	
2	2	1	NaN	NaN	NaN	NaN	5.0	
3	2	2	NaN	NaN	NaN	NaN	6.0	
4	2	3	NaN	NaN	NaN	NaN	6.0	

	pregend2	nbrnaliv	multbrth	...	laborfor_i	religion_i	metro_i	\
0	NaN	1.0	NaN	...	0	0	0	
1	NaN	1.0	NaN	...	0	0	0	
2	NaN	3.0	5.0	...	0	0	0	
3	NaN	1.0	NaN	...	0	0	0	
4	NaN	1.0	NaN	...	0	0	0	

	basewgt	adj_mod_basewgt	finalwgt	secu_p	sest	cmintvw	\
0	3410.389399	3869.349602	6448.271112	2	9	NaN	
1	3410.389399	3869.349602	6448.271112	2	9	NaN	
2	7226.301740	8567.549110	12999.542264	2	12	NaN	
3	7226.301740	8567.549110	12999.542264	2	12	NaN	

```
4    7226.301740      8567.549110  12999.542264      2    12      NaN
```

```
    totalwgt_lb
0         8.8125
1         7.8750
2         9.1250
3         7.0000
4         6.1875
```

[5 rows x 244 columns]

Print the column names.

```
[7]: preg.columns
```

```
[7]: Index(['caseid', 'pregordr', 'howpreg_n', 'howpreg_p', 'moscurrp', 'nowprgdk',
          'pregend1', 'pregend2', 'nbrnaliv', 'multbrth',
          ...,
          'laborfor_i', 'religion_i', 'metro_i', 'basewgt', 'adj_mod_basewgt',
          'finalwgt', 'secu_p', 'sest', 'cmintvw', 'totalwgt_lb'],
          dtype='object', length=244)
```

Select a single column name.

```
[9]: preg.columns[1]
```

```
[9]: 'pregordr'
```

Select a column and check what type it is.

```
[10]: pregordr = preg['pregordr']
      type(pregordr)
```

```
[10]: pandas.core.series.Series
```

Print a column.

```
[11]: pregordr
```

```
[11]: 0         1
      1         2
      2         1
      3         2
      4         3
      ..
     13588      1
     13589      2
     13590      3
     13591      4
     13592      5
```

Name: pregordr, Length: 13593, dtype: int64

Select a single element from a column.

```
[12]: pregordr[0]
```

```
[12]: 1
```

Select a slice from a column.

```
[13]: pregordr[2:5]
```

```
[13]: 2    1
      3    2
      4    3
```

Name: pregordr, dtype: int64

Select a column using dot notation.

```
[84]: pregordr = preg.pregordr
      pregordr.head()
```

```
[84]: 0    1
      1    2
      2    1
      3    2
      4    3
```

Name: pregordr, dtype: int64

Count the number of times each value occurs.

```
[15]: caseid = 10229
      preg_map = nsfg.MakePregMap(preg)
      indices = preg_map[caseid]
      preg.outcome[indices].value_counts().sort_index()
```

```
[15]: 1    9148
      2    1862
      3     120
      4    1921
      5     190
      6     352
```

Name: outcome, dtype: int64

Check the values of another variable.

```
[16]: preg.birthwgt_lb.value_counts().sort_index()
```

```
[16]: 0.0      8
      1.0     40
```

```

2.0      53
3.0      98
4.0     229
5.0     697
6.0    2223
7.0    3049
8.0    1889
9.0     623
10.0    132
11.0     26
12.0     10
13.0      3
14.0      3
15.0      1
Name: birthwgt_lb, dtype: int64

```

Make a dictionary that maps from each respondent's `caseid` to a list of indices into the pregnancy DataFrame. Use it to select the pregnancy outcomes for a single respondent.

```

[17]: caseid = 10229
      preg_map = nsfg.MakePregMap(preg)
      indices = preg_map[caseid]
      preg.outcome[indices].values

```

```

[17]: array([4, 4, 4, 4, 4, 4, 1], dtype=int64)

```

1.2 Exercises

Select the `birthord` column, print the value counts, and compare to results published in the [codebook](#)

```

[56]: birthord = preg['birthord']
      type(birthord)
      birthord
      preg.birthord.value_counts().sort_index()

```

```

[56]: 1.0      4413
      2.0     2874
      3.0     1234
      4.0      421
      5.0     126
      6.0      50
      7.0      20
      8.0       7
      9.0       2
      10.0      1
      Name: birthord, dtype: int64

```

We can also use `isnull` to count the number of nans.

```
[19]: preg.birthord.isnull().sum()
```

```
[19]: 4445
```

Select the `prglngth` column, print the value counts, and compare to results published in the [codebook](#)

```
[20]: prglngth = preg['prglngth']
type(prglngth)
prglngth
preg.prglngth.value_counts().sort_index()
```

```
[20]: 0      15
      1      9
      2     78
      3    151
      4    412
      5    181
      6    543
      7    175
      8    409
      9    594
     10    137
     11    202
     12    170
     13    446
     14     29
     15     39
     16     44
     17    253
     18     17
     19     34
     20     18
     21     37
     22    147
     23     12
     24     31
     25     15
     26    117
     27      8
     28     38
     29     23
     30    198
     31     29
     32    122
     33     50
     34     60
     35    357
```

```

36      329
37      457
38      609
39     4744
40     1120
41      591
42      328
43      148
44       46
45       10
46        1
47        1
48        7
50        2

```

Name: prglngth, dtype: int64

To compute the mean of a column, you can invoke the `mean` method on a Series. For example, here is the mean birthweight in pounds:

```
[21]: preg.totalwgt_lb.mean()
```

```
[21]: 7.265628457623368
```

Create a new column named `totalwgt_kg` that contains birth weight in kilograms. Compute its mean. Remember that when you create a new column, you have to use dictionary syntax, not dot notation.

```
[34]: preg['totalwgt_kg']=preg.totalwgt_lb * 0.45359237
preg.totalwgt_kg
preg.columns
```

```
[34]: Index(['caseid', 'pregordr', 'howpreg_n', 'howpreg_p', 'moscurrp', 'nowprgdk',
          'pregend1', 'pregend2', 'nbrnaliv', 'multbrth',
          ...,
          'religion_i', 'metro_i', 'basewgt', 'adj_mod_basewgt', 'finalwgt',
          'secu_p', 'sest', 'cmintvw', 'totalwgt_lb', 'totalwgt_kg'],
          dtype='object', length=245)
```

`nsfg.py` also provides `ReadFemResp`, which reads the female respondents file and returns a `DataFrame`:

```
[38]: resp = nsfg.ReadFemResp()
```

`DataFrame` provides a method `head` that displays the first five rows:

```
[37]: resp.head()
```

```
[37]:   caseid  rscrinf  rdormres  rostscrn  rscreenhisp  rscreenrace  age_a  \
0    2298        1         5         5           1          5.0    27
```

1	5012	1	5	1	5	5.0	42
2	11586	1	5	1	5	5.0	43
3	6794	5	5	4	1	5.0	15
4	616	1	5	4	1	5.0	20

	age_r	cmbirth	agescrn	...	pubassis_i	basewgt	adj_mod_basewgt	\
0	27	902	27	...	0	3247.916977	5123.759559	
1	42	718	42	...	0	2335.279149	2846.799490	
2	43	708	43	...	0	2335.279149	2846.799490	
3	15	1042	15	...	0	3783.152221	5071.464231	
4	20	991	20	...	0	5341.329968	6437.335772	

	finalwgt	secu_r	sest	cmintvw	cmlstyr	screentime	intvlngh
0	5556.717241	2	18	1234	1222	18:26:36	110.492667
1	4744.191350	2	18	1233	1221	16:30:59	64.294000
2	4744.191350	2	18	1234	1222	18:19:09	75.149167
3	5923.977368	2	18	1234	1222	15:54:43	28.642833
4	7229.128072	2	18	1233	1221	14:19:44	69.502667

[5 rows x 3087 columns]

Select the `age_r` column from `resp` and print the value counts. How old are the youngest and oldest respondents?

```
[42]: age_r = resp['age_r']
type(age_r)
age_r
resp.age_r.value_counts().sort_index()
# Youngest is 15 years old and the oldest is 44 years old based on the results
→ shown below.
```

```
[42]: 15    217
16    223
17    234
18    235
19    241
20    258
21    267
22    287
23    282
24    269
25    267
26    260
27    255
28    252
29    262
30    292
31    278
```

```

32    273
33    257
34    255
35    262
36    266
37    271
38    256
39    215
40    256
41    250
42    215
43    253
44    235

```

Name: age_r, dtype: int64

We can use the `caseid` to match up rows from `resp` and `preg`. For example, we can select the row from `resp` for `caseid` 2298 like this:

```
[43]: resp[resp.caseid==2298]
```

```

[43]:   caseid  rscrinf  rdormres  rostscrn  rscreenhisp  rscreenrace  age_a  \
0    2298         1         5         5         1         5.0    27

      age_r  cmbirth  agescrn  ...  pubassis_i      basewgt  adj_mod_basewgt  \
0     27      902      27  ...          0  3247.916977    5123.759559

      finalwgt  secu_r  sest  cmintvw  cmlstyr  screentime  intvlngh
0  5556.717241      2   18    1234    1222    18:26:36  110.492667

[1 rows x 3087 columns]

```

And we can get the corresponding rows from `preg` like this:

```
[65]: preg[preg.caseid==2298]
```

```

[65]:   caseid  pregordr  howpreg_n  howpreg_p  moscurrp  nowprgdk  pregend1  \
2610    2298         1        NaN        NaN        NaN        NaN        6.0
2611    2298         2        NaN        NaN        NaN        NaN        6.0
2612    2298         3        NaN        NaN        NaN        NaN        6.0
2613    2298         4        NaN        NaN        NaN        NaN        6.0

      pregend2  nbrnaliv  multbrth  ...  laborfor_i  religion_i  metro_i  \
2610      NaN      1.0      NaN  ...          0          0          0
2611      NaN      1.0      NaN  ...          0          0          0
2612      NaN      1.0      NaN  ...          0          0          0
2613      NaN      1.0      NaN  ...          0          0          0

      basewgt  adj_mod_basewgt      finalwgt  secu_p  sest  cmintvw  \

```


2610	3247.916977	5123.759559	5556.717241	2	18	NaN
2611	3247.916977	5123.759559	5556.717241	2	18	NaN
2612	3247.916977	5123.759559	5556.717241	2	18	NaN
2613	3247.916977	5123.759559	5556.717241	2	18	NaN

totalwgt_lb	
2610	6.8750
2611	5.5000
2612	4.1875
2613	6.8750

[4 rows x 244 columns]

How old is the respondent with caseid 1?

```
[46]: resp[resp.caseid==1]
# The respondent with caseid 1 is 44 years old as the results show below.
```

```
[46]:      caseid  rscrinf  rdormres  rostscrn  rscreenhisp  rscreenrace  age_a \
1069      1      1      5      4      5      5.0      44

      age_r  cmbirth  agescrn  ...  pubassis_i      basewgt  adj_mod_basewgt  \
1069     44      695      44  ...      0  3410.389399      3869.349602

      finalwgt  secu_r  sest  cmintvw  cmlstyr  screentime  intvlngth
1069  6448.271112      2    9    1231    1219    19:56:43  67.563833
```

[1 rows x 3087 columns]

What are the pregnancy lengths for the respondent with caseid 2298?

```
[75]: preg.shape
preg.dtypes
preg.info()

preg2=preg[["caseid","prglngth"]]
preg2

preg2[preg2.caseid==2298]
#There were four pregnancies related to caseid 2298. Lengths were 40, 36, 30,
↳ and 40 weeks.
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 13593 entries, 0 to 13592
Columns: 244 entries, caseid to totalwgt_lb
dtypes: float64(171), int64(73)
memory usage: 25.3 MB
```

```
[75]:      caseid  prglngth
      2610    2298      40
      2611    2298      36
      2612    2298      30
      2613    2298      40
```

What was the birthweight of the first baby born to the respondent with caseid 5012?

```
[81]: preg3=preg[["caseid","birthord","birthwgt_lb"]]
      preg3

      preg3[preg3.caseid==5012]
      #Birthweight of the first baby born was 6 lbs for respondent with caseid 5012.
```

```
[81]:      caseid  birthord  birthwgt_lb
      5515    5012      1.0      6.0
```