

# chap02ex

December 19, 2021

## 1 Examples and Exercises from Think Stats, 2nd Edition

<http://thinkstats2.com>

Copyright 2016 Allen B. Downey

MIT License: <https://opensource.org/licenses/MIT>

By Saima Rahmanzai DSC530 Chapter02 Exercise

```
[1]: from __future__ import print_function, division

    %matplotlib inline

    import numpy as np
    import pandas as pd

    import nsfg
    import first
```

Given a list of values, there are several ways to count the frequency of each value.

```
[2]: t = [1, 2, 2, 3, 5]
```

You can use a Python dictionary:

```
[3]: hist = {}
    for x in t:
        hist[x] = hist.get(x, 0) + 1

    hist
```

```
[3]: {1: 1, 2: 2, 3: 1, 5: 1}
```

You can use a `Counter` (which is a dictionary with additional methods):

```
[4]: from collections import Counter
    counter = Counter(t)
    counter
```

```
[4]: Counter({1: 1, 2: 2, 3: 1, 5: 1})
```

Or you can use the `Hist` object provided by `thinkstats2`:

```
[5]: import thinkstats2
hist = thinkstats2.Hist([1, 2, 2, 3, 5])
hist
```

```
[5]: Hist({1: 1, 2: 2, 3: 1, 5: 1})
```

`Hist` provides `Freq`, which looks up the frequency of a value.

```
[6]: hist.Freq(2)
```

```
[6]: 2
```

You can also use the bracket operator, which does the same thing.

```
[7]: hist[2]
```

```
[7]: 2
```

If the value does not appear, it has frequency 0.

```
[8]: hist[4]
```

```
[8]: 0
```

The `Values` method returns the values:

```
[9]: hist.Values()
```

```
[9]: dict_keys([1, 2, 3, 5])
```

So you can iterate the values and their frequencies like this:

```
[10]: for val in sorted(hist.Values()):
        print(val, hist[val])
```

```
1 1
2 2
3 1
5 1
```

Or you can use the `Items` method:

```
[11]: for val, freq in hist.Items():
        print(val, freq)
```


```
1 1
2 2
3 1
5 1
```

thinkplot is a wrapper for matplotlib that provides functions that work with the objects in thinkstats2.

For example Hist plots the values and their frequencies as a bar graph.

Config takes parameters that label the x and y axes, among other things.

```
[12]: import thinkplot
      thinkplot.Hist(hist)
      thinkplot.Config(xlabel='value', ylabel='frequency')
```




As an example, I'll replicate some of the figures from the book.

First, I'll load the data from the pregnancy file and select the records for live births.

```
[13]: preg = nsfg.ReadFemPreg()
      live = preg[preg.outcome == 1]
```

Here's the histogram of birth weights in pounds. Notice that Hist works with anything iterable, including a Pandas Series. The label attribute appears in the legend when you plot the Hist.


```
[14]: hist = thinkstats2.Hist(live.birthwgt_lb, label='birthwgt_lb')
      thinkplot.Hist(hist)
      thinkplot.Config(xlabel='Birth weight (pounds)', ylabel='Count')
```



Before plotting the ages, I'll apply `floor` to round down:

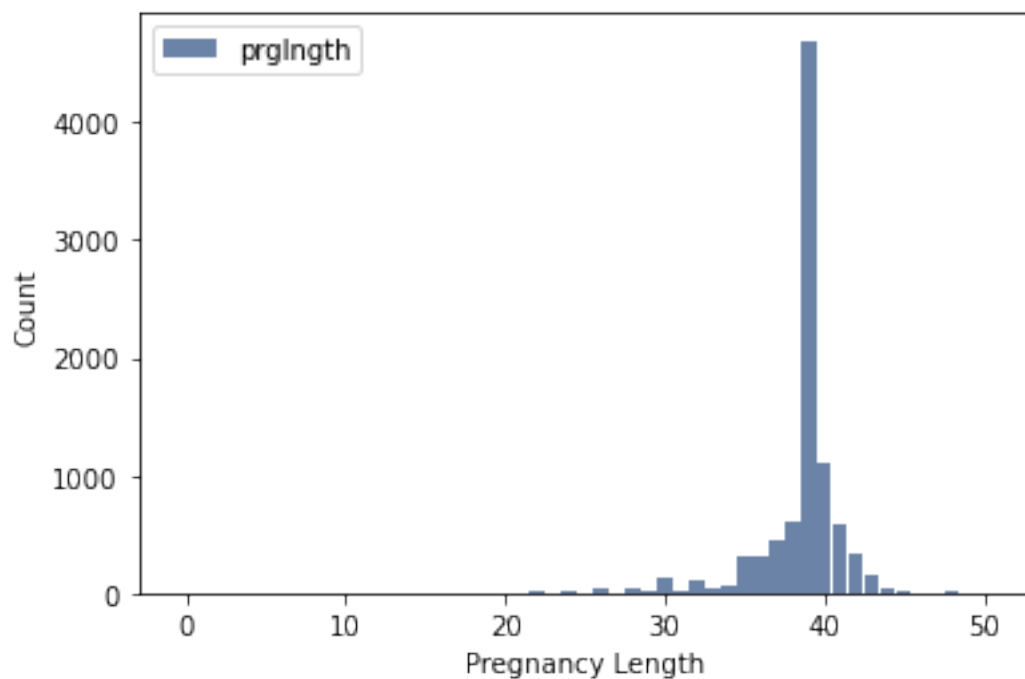
```
[20]: ages = np.floor(live.agepreg)
```

```
[21]: hist = thinkstats2.Hist(ages, label='agepreg')
      thinkplot.Hist(hist)
      thinkplot.Config(xlabel='years', ylabel='Count')
```



As an exercise, plot the histogram of pregnancy lengths (column `prglength`).

```
[23]: hist = thinkstats2.Hist(live.prglength, label='prglength')
      thinkplot.Hist(hist)
      thinkplot.Config(xlabel='Pregnancy Length', ylabel='Count')
```



Hist provides `smallest`, which select the lowest values and their frequencies.

```
[24]: for weeks, freq in hist.Smallest(10):  
      print(weeks, freq)
```

```
0 1  
4 1  
9 1  
13 1  
17 2  
18 1  
19 1  
20 1  
21 2  
22 7
```

Use `Largest` to display the longest pregnancy lengths.

```
[26]: for weeks, freq in hist.Largest(10):  
      print(weeks, freq)
```


```
50 2  
48 7  
47 1  
46 1  
45 10  
44 46  
43 148  
42 328  
41 587  
40 1116
```

From live births, we can select first babies and others using `birthord`, then compute histograms of pregnancy length for the two groups.

```
[56]: firsts = live[live.birthord == 1]  
      others = live[live.birthord != 1]  
  
      first_hist = thinkstats2.Hist(firsts.prglength, label='first')  
      other_hist = thinkstats2.Hist(others.prglength, label='other')
```

We can use `width` and `align` to plot two histograms side-by-side.

```
[31]: width = 0.45  
      thinkplot.PrePlot(2)  
      thinkplot.Hist(first_hist, align='right', width=width)  
      thinkplot.Hist(other_hist, align='left', width=width)  
      thinkplot.Config(xlabel='weeks', ylabel='Count', xlim=[27, 46])
```



Series provides methods to compute summary statistics:

```
[32]: mean = live.prglngh.mean()
      var = live.prglngh.var()
      std = live.prglngh.std()
```

Here are the mean and standard deviation:

```
[33]: mean, std
```

```
[33]: (38.56055968517709, 2.702343810070593)
```

As an exercise, confirm that `std` is the square root of `var`:

```
[36]: import math
      math.sqrt(var)
      #As seen above, I used the math function to perform a square root on the
      ↪ variance already calculated.
      #It came to the same value as the standard deviation of 2.702 which confirms
      ↪ that std is the
      #square root of variance.
```

```
[36]: 2.702343810070593
```

Here's are the mean pregnancy lengths for first babies and others:

```
[37]: firsts.prglngh.mean(), others.prglngh.mean()
```

```
[37]: (38.60095173351461, 38.52291446673706)
```

And here's the difference (in weeks):

```
[38]: firsts.prglength.mean() - others.prglength.mean()
```

```
[38]: 0.07803726677754952
```

This function computes the Cohen effect size, which is the difference in means expressed in number of standard deviations:

```
[41]: def CohenEffectSize(group1, group2):  
    """Computes Cohen's effect size for two groups.  
  
    group1: Series or DataFrame  
    group2: Series or DataFrame  
  
    returns: float if the arguments are Series;  
            Series if the arguments are DataFrames  
    """  
    diff = group1.mean() - group2.mean()  
  
    var1 = group1.var()  
    var2 = group2.var()  
    n1, n2 = len(group1), len(group2)  
  
    pooled_var = (n1 * var1 + n2 * var2) / (n1 + n2)  
    d = diff / np.sqrt(pooled_var)  
    return d
```

Compute the Cohen effect size for the difference in pregnancy length for first babies and others.

```
[73]: group1 = live[live.birthord == 1] #Firsts  
    group2 = live[live.birthord != 1] #Others  
  
    diff = group1.prglength.mean() - group2.prglength.mean() #diff is 0.078  
  
    var1 = group1.prglength.var()  
    var2 = group2.prglength.var()  
    n1, n2 = len(group1), len(group2)  
  
    pooled_var = (n1 * var1 + n2 * var2) / (n1 + n2)  
    d = diff / np.sqrt(pooled_var)  
  
    print(d)
```

```
0.028879044654449883
```



## 1.1 Exercises

Using the variable `totalwgt_lb`, investigate whether first babies are lighter or heavier than others.

Compute Cohen's effect size to quantify the difference between the groups. How does it compare to the difference in pregnancy length?

```
[65]: group1 = live[live.birthord == 1] #Firsts
      group2 = live[live.birthord != 1] #Others

      diff2 = group1.totalwgt_lb.mean() - group2.totalwgt_lb.mean()
      diff2
      #On an aggregate, the mean of totalwgt_lbs is smaller than that of others but
      ↪only slightly.
```

```
[65]: -0.12476118453549034
```

```
[71]: var1 = group1.totalwgt_lb.var()
      var2 = group2.totalwgt_lb.var()
      n1, n2 = len(group1), len(group2)

      n1 #There are 4413 firsts
      n2 #There are 4735 others

      pooled_var = (n1 * var1 + n2 * var2) / (n1 + n2)
      d2 = diff / np.sqrt(pooled_var)

      print(d2)

      # The Cohen's difference is not very significant 0.05 especially when the firsts
      ↪numbers are also smaller
      # than the others group that could effect the results.
```


```
0.055464308804653806
```

```
[80]: # import matplotlib.pyplot as plt
      # %matplotlib inline

      firsts = live[live.birthord == 1] #Firsts
      others = live[live.birthord != 1] #Others
      first_hist = thinkstats2.Hist(firsts.totalwgt_lb, label='first')
      other_hist = thinkstats2.Hist(others.totalwgt_lb, label='other')

      width = 0.45
      thinkplot.PrePlot(2)
      thinkplot.Hist(first_hist, align='right', width=width)
      thinkplot.Hist(other_hist, align='left', width=width)
```

```
thinkplot.Config(xlabel='weight', ylabel='Count')
```



For the next few exercises, we'll load the respondent file:

```
[35]: #Understanding teh ReadFemResp and the variable totincr prior to creating the
      ↪ histogram
resp = nsfg.ReadFemResp()
resp.head()
resp.columns
resp.info
resp2=resp[["caseid","totincr"]]
resp2.head()
resp2.totincr.value_counts().sort_index()
```

```
[35]: 1      299
      2      301
      3      266
      4      421
      5      445
      6      559
      7      583
      8      606
      9      607
     10      468
     11      647
```

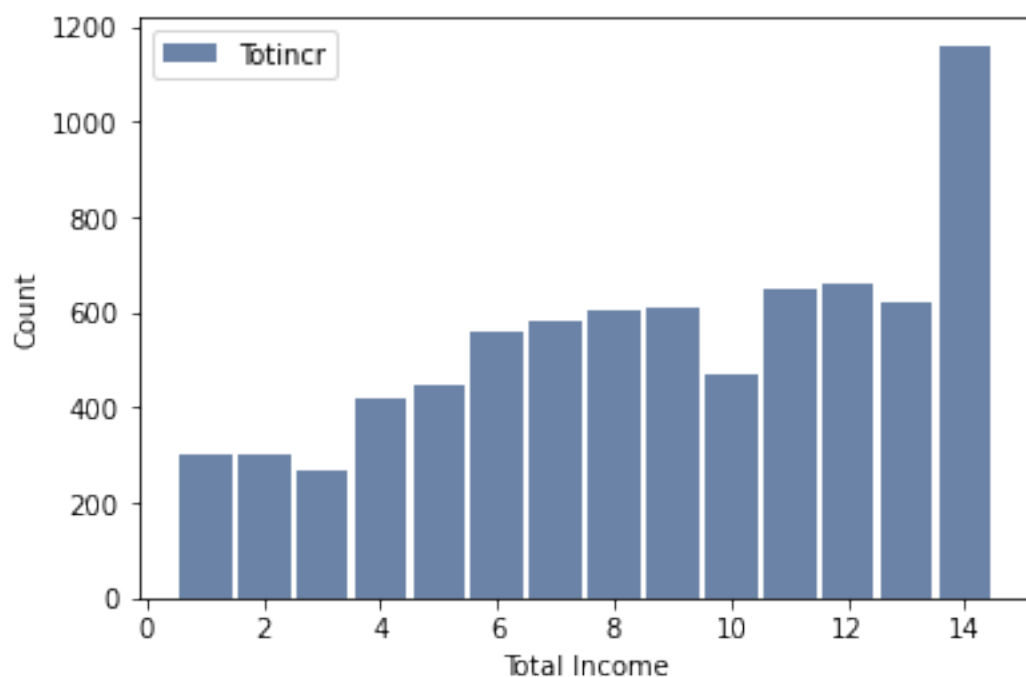
```
12      658
13      623
14     1160
Name: totincr, dtype: int64
```

```
[36]: resp2.totincr.value_counts().sum()
#per the codebook, total value label total was 7643 and was validated with the
↪data with the command
#above.
```

```
[36]: 7643
```

Make a histogram of totincr the total income for the respondent's family. To interpret the codes see the [codebook](#).

```
[37]: hist = thinkstats2.Hist(resp2.totincr, label='Totincr')
thinkplot.Hist(hist)
thinkplot.Config(xlabel='Total Income', ylabel='Count')
```



Make a histogram of age\_r, the respondent's age at the time of interview.

```
[38]: resp = nsfg.ReadFemResp()
resp.head()
resp3=resp[["caseid", "age_r"]]
resp3.head()
```


```
resp3.age_r.value_counts().sort_index()
```

```
[38]: 15    217
      16    223
      17    234
      18    235
      19    241
      20    258
      21    267
      22    287
      23    282
      24    269
      25    267
      26    260
      27    255
      28    252
      29    262
      30    292
      31    278
      32    273
      33    257
      34    255
      35    262
      36    266
      37    271
      38    256
      39    215
      40    256
      41    250
      42    215
      43    253
      44    235
      Name: age_r, dtype: int64
```

```
[40]: resp3.age_r.value_counts().sum()
```

```
[40]: 7643
```

```
[41]: # Histogram for age_r
      hist = thinkstats2.Hist(resp3.age_r, label='Age')
      thinkplot.Hist(hist)
      thinkplot.Config(xlabel='Age', ylabel='Count')
```



Make a histogram of numfmhh, the number of people in the respondent's household.

```
[42]: resp = nsfg.ReadFemResp()
      resp.head()
      resp4=resp[["caseid", "numfmhh"]]
      resp4.head()
      resp4.numfmhh.value_counts().sort_index()
```


```
[42]: 0      942
      1     1716
      2     1826
      3     1740
      4      906
      5      313
      6      118
      7       78
      8         4
      Name: numfmhh, dtype: int64
```

```
[43]: resp4.numfmhh.value_counts().sum()
```

```
[43]: 7643
```

```
[44]: # Histogram for numfmhh
      hist = thinkstats2.Hist(resp4.numfmhh, label='No. of People')
```

```
thinkplot.Hist(hist)
thinkplot.Config(xlabel='No. of People', ylabel='Count')
```



Make a histogram of parity, the number of children borne by the respondent. How would you describe this distribution?

```
[46]: resp = nsfg.ReadFemResp()
      resp.head()
      resp5=resp[["caseid","parity"]]
      resp5.head()
      resp5.parity.value_counts().sort_index()
```


```
[46]: 0    3230
      1    1519
      2    1603
      3     828
      4     309
      5      95
      6      29
      7      15
      8       8
      9       2
     10       3
     16       1
     22       1
```

Name: parity, dtype: int64

```
[47]: resp5.parity.value_counts().sum()
```

```
[47]: 7643
```

```
[48]: # Histogram for parity
hist = thinkstats2.Hist(resp5.parity, label='Parity')
thinkplot.Hist(hist)
thinkplot.Config(xlabel='Parity', ylabel='Count')
#Parity is defined as the number of times that a respondent has given birth to
→a fetus with a gestational
#age of 24 weeks or more. The graph shows 0 babies to respondents being the
→highest value. More
#testing below.
```



Use Hist.Largest to find the largest values of parity.

```
[50]: import thinkstats2
hist = thinkstats2.Hist(resp5.parity)
hist
```

```
[50]: Hist({0: 3230, 2: 1603, 1: 1519, 3: 828, 4: 309, 5: 95, 6: 29, 7: 15, 8: 8, 10:
3, 9: 2, 22: 1, 16: 1})
```

```
[69]: for parity, freq in hist.Largest(10):
      print(parity, freq)
      # or
      group1.parity.value_counts().sort_index()
```

```
8 1
7 1
5 5
4 19
3 123
2 267
1 229
0 515
```

```
[69]: 0    515
      1    229
      2    267
      3    123
      4     19
      5      5
      7      1
      8      1
      Name: parity, dtype: int64
```

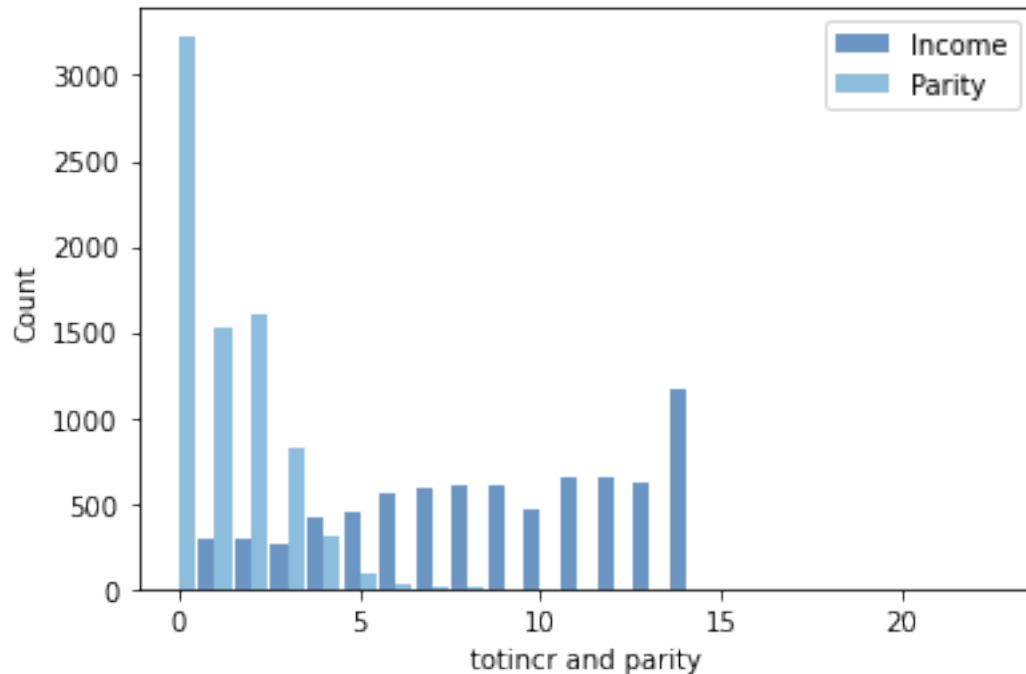
Let's investigate whether people with higher income have higher parity. Keep in mind that in this study, we are observing different people at different times during their lives, so this data is not the best choice for answering this question. But for now let's take it at face value.

Use `totincr` to select the respondents with the highest income (level 14). Plot the histogram of parity for just the high income respondents.

```
[54]: Income = resp2.totincr
      Parity = resp5.parity
      Income_hist = thinkstats2.Hist(Income, label='Income')
      Parity_hist = thinkstats2.Hist(Parity, label='Parity')

      width = 0.45
      thinkplot.PrePlot(2)
      thinkplot.Hist(Income_hist, align='right', width=width)
      thinkplot.Hist(Parity_hist, align='left', width=width)
      thinkplot.Config(xlabel='totincr and parity', ylabel='Count')
      #As shown by the histogram below, the incomee and parity are on the opposite_
      ↪ends. High parity
      #showing 0 babies appear to be corresponding with low income group.
```





Find the largest parities for high income respondents.

```
[68]: group1 = resp6[resp6.totincr == 14] #High income

import thinkstats2
hist = thinkstats2.Hist(group1.parity)
hist

for parity, freq in hist.Largest(10):
    print(parity, freq)
```

```
8 1
7 1
5 5
4 19
3 123
2 267
1 229
0 515
```

Compare the mean parity for high income respondents and others.

```
[60]: group1 = resp6[resp6.totincr == 14] #High income
group2 = resp6[resp6.totincr != 14] #Others

diff3 = group1.parity.mean() - group2.parity.mean()
```

```
diff3
```

```
[60]: -0.17371374470099532
```

Compute the Cohen effect size for this difference. How does it compare with the difference in pregnancy length for first babies and others?

```
[65]: var1 = group1.parity.var()
var2 = group2.parity.var()
n1, n2 = len(group1), len(group2)

n1 #There are 1160 records with high income
n2 #There are 6483 records with not high income

pooled_var3 = (n1 * var1 + n2 * var2) / (n1 + n2)
d3 = diff3 / np.sqrt(pooled_var3)

print(d3)
```

```
-0.1251185531466061
```

```
[62]: group1 = resp6[resp6.totincr == 14] #High income
group2 = resp6[resp6.totincr != 14] #Others
first_hist = thinkstats2.Hist(group1.parity, label='high income 14')
other_hist = thinkstats2.Hist(group2.parity, label='Not high income')

width = 0.45
thinkplot.PrePlot(2)
thinkplot.Hist(first_hist, align='right', width=width)
thinkplot.Hist(other_hist, align='left', width=width)
thinkplot.Config(xlabel='Parity by income', ylabel='Count')
```

